Chapter 4 Linear Models for Classification

Yue Yu

4.1 Discriminant Functions

Skipped reading.

4.2 Probabilistic Generative Models

(4.57)

$$\begin{aligned} p(\mathcal{C}_1|\mathbf{x}) &= \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x})} \\ &= \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{\sum_{k=1}^{K} p(\mathbf{x}, \mathcal{C}_k)} \\ &= \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1) + p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)} \\ &= \frac{1}{1 + \frac{p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)}{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}} \\ &= \frac{1}{1 + \exp(-a)}. \end{aligned}$$

(4.65) – (4.67)

We can readily derive (4.65) by noticing that all terms will be canceled out except for those containing $\boldsymbol{\mu}_k$, provided (4.66) and (4.67).

(4.73)

As given by (4.72), the terms in the log likelihood depending on π are

$$\sum_{n=1}^{N} \{t_n \ln \pi + (1 - t_n) \ln(1 - \pi)\}.$$

Setting the derivative of the log likelihood function with respect to π to 0, we have

$$\begin{aligned} \frac{\partial}{\partial \pi} \ell(\pi, \boldsymbol{\mu}_1, \boldsymbol{\mu}_2, \boldsymbol{\Sigma}) &= \sum_{n=1}^{N} t_n \frac{1}{\pi} - \sum_{n=1}^{N} (1 - t_n) \frac{1}{1 - \pi} \\ &= 0. \end{aligned}$$

Solving for π while denoting the total number of data points in class $\mathcal{C}_1$ by N_1, we obtain

$$\pi = \frac{N_1}{N},$$

which is the fraction of points in class $\mathcal{C}_1$.

This can be generalized to $K > 2$ classes where $\mathbf{t}_n$ is a one hot vector of length K such that $t_{nj} = I_{jk}$. Then, the likelihood function can be written as

$$p(\mathbf{X}, \mathbf{T}|\pi_k, \boldsymbol{\mu}_k, \boldsymbol{\Sigma}) = \prod_{n=1}^{N} \prod_{k=1}^{K} (\pi_k \mathcal{N}(\mathbf{x}_n|\boldsymbol{\mu}_k, \boldsymbol{\Sigma}))^{t_{nk}}.$$

The corresponding log likelihood function is

$$\ell(\pi_k, \boldsymbol{\mu}_k, \boldsymbol{\Sigma}) = \sum_{n=1}^{N} \sum_{k=1}^{K} t_{nk} (\ln \pi_k + \ln \mathcal{N}(\mathbf{x}_n|\boldsymbol{\mu}_k, \boldsymbol{\Sigma})).$$

Here, we are only interested in the terms depending on π_k, namely,

$$\sum_{n=1}^{N} \sum_{k=1}^{K} t_{nk} \ln \pi_k.$$

To find π_k, we construct the Lagrangian using the constraint $\sum_{k=1}^{K} \pi_k = 1$, given by

$$\mathcal{L}(\pi_k, \lambda) = \sum_{n=1}^{N} \sum_{k=1}^{K} t_{nk} \ln \pi_k + \lambda \Big(\sum_{k=1}^{K} \pi_k - 1 \Big).$$

Setting the derivative with respect to π_k to 0, we have

$$\begin{aligned} \frac{\partial}{\partial \pi_k} \mathcal{L}(\pi_k, \lambda) &= \sum_{n=1}^{N} t_{nk} \frac{1}{\pi_k} + \lambda \\ &= 0. \end{aligned}$$

Solving for π_k, we obtain

$$\pi_k = -\frac{1}{\lambda} \sum_{n=1}^{N} t_{nk} = -\frac{1}{\lambda} N_k. \tag{*}$$

Summing over k on both sides, we have

$$\sum_{k=1}^{K} \pi_k = -\frac{N}{\lambda} = 1,$$

which implies that

$$\lambda = -N.$$

Substituting back into (*), we obtain

$$\pi_k = \frac{N_k}{N},$$

which is the fraction of points in class $\mathcal{C}_k$.

(4.75) – (4.76)

To find $\boldsymbol{\mu}_1$, we set the derivative of the log likelihood with respect to $\boldsymbol{\mu}_1$ to $\mathbf{0}$,

$$\begin{aligned} \frac{\partial}{\partial \boldsymbol{\mu}_1} \ell(\pi, \boldsymbol{\mu}_1, \boldsymbol{\mu}_2, \boldsymbol{\Sigma}) &= \frac{\partial}{\partial \boldsymbol{\mu}_1} \sum_{n=1}^{N} t_n \ln \mathcal{N}(\mathbf{x}_n|\boldsymbol{\mu}_1, \boldsymbol{\Sigma}) \\ &= -\frac{1}{2} \frac{\partial}{\partial \boldsymbol{\mu}_1} \sum_{n=1}^{N} t_n (\mathbf{x}_n - \boldsymbol{\mu}_1)^{\mathrm{T}} \boldsymbol{\Sigma}^{-1} (\mathbf{x}_n - \boldsymbol{\mu}_1) \\ &= \sum_{n=1}^{N} t_n (-\boldsymbol{\Sigma}^{-1} \boldsymbol{\mu}_1 + \boldsymbol{\Sigma}^{-1} \mathbf{x}_n) \\ &= \mathbf{0}. \end{aligned}$$

Solving for $\boldsymbol{\mu}_1$, we obtain

$$\boldsymbol{\mu}_1 = \frac{1}{N_1}\sum_{n=1}^{N} t_n \mathbf{x}_n,$$

where we denote $N_1 = \sum_{n=1}^{N} t_n$ as the number of data points assigned to class $\mathcal{C}_1$. Similarly,

$$\boldsymbol{\mu}_2 = \frac{1}{N_2}\sum_{n=1}^{N} (1-t_n) \mathbf{x}_n,$$

where we denote $N_2 = \sum_{n=1}^{N} (1-t_n)$ as the number of data points assigned to class $\mathcal{C}_2$.

This can be generalized to $K > 2$ classes with the same settings as the derivation of (4.73). The log likelihood function is

$$\ell(\pi_k, \boldsymbol{\mu}_k, \boldsymbol{\Sigma}) = \sum_{n=1}^{N}\sum_{k=1}^{K} t_{nk}(\ln \pi_k + \ln \mathcal{N}(\mathbf{x}_n|\boldsymbol{\mu}_k, \boldsymbol{\Sigma})).$$

Here we are only interested in $\boldsymbol{\mu}_k$. Setting the derivative with respect to $\boldsymbol{\mu}_k$,

$$\begin{aligned}
\frac{\partial}{\partial \boldsymbol{\mu}_k}\ell(\pi_k, \boldsymbol{\mu}_k, \boldsymbol{\Sigma}) &= -\frac{1}{2}\frac{\partial}{\partial \boldsymbol{\mu}_k}\sum_{n=1}^{N}\sum_{k=1}^{K} t_{nk}(\mathbf{x}_n - \boldsymbol{\mu}_k)^{\mathrm{T}}\boldsymbol{\Sigma}^{-1}(\mathbf{x} - \boldsymbol{\mu}_k) \\
&= \sum_{n=1}^{N} t_{nk}(-\boldsymbol{\Sigma}^{-1}\boldsymbol{\mu}_k + \boldsymbol{\Sigma}^{-1}\mathbf{x}_n) \\
&= \mathbf{0}.
\end{aligned}$$

Solving for $\boldsymbol{\mu}_k$, we obtain

$$\boldsymbol{\mu}_k = \frac{1}{N_k}\sum_{n=1}^{N} t_{nk}\mathbf{x}_n,$$

where $N_k = \sum_{n=1}^{N} t_{nk}$, representing the number of data points that are assigned to class $\mathcal{C}_k$.

(4.77) – (4.80)

To find $\boldsymbol{\Sigma}$, we set the derivative of the log likelihood function with respect to $\boldsymbol{\Sigma}^{-1}$ to $\mathbf{0}$. Specifically,

$$\begin{aligned}
\frac{\partial}{\partial \boldsymbol{\Sigma}^{-1}}\ell(\pi, \boldsymbol{\mu}_1, \boldsymbol{\mu}_2, \boldsymbol{\Sigma}) &= \frac{\partial}{\partial \boldsymbol{\Sigma}^{-1}}\bigg(-\frac{1}{2}\sum_{n=1}^{N} t_n \ln|\boldsymbol{\Sigma}| - \frac{1}{2}\sum_{n=1}^{N} t_n(\mathbf{x}_n - \boldsymbol{\mu}_1)^{\mathrm{T}}\boldsymbol{\Sigma}^{-1}(\mathbf{x}_n - \boldsymbol{\mu}_1) \\
&\qquad - \frac{1}{2}\sum_{n=1}^{N}(1-t_n)\ln|\boldsymbol{\Sigma}| - \frac{1}{2}\sum_{n=1}^{N}(1-t_n)(\mathbf{x}_n - \boldsymbol{\mu}_2)^{\mathrm{T}}\boldsymbol{\Sigma}^{-1}(\mathbf{x}_n - \boldsymbol{\mu}_2)\bigg) \\
&= \frac{\partial}{\partial \boldsymbol{\Sigma}^{-1}}\bigg(-\frac{N}{2}\ln|\boldsymbol{\Sigma}| - \frac{1}{2}\sum_{n=1}^{N} t_n \mathrm{Tr}\big\{(\mathbf{x}_n - \boldsymbol{\mu}_1)^{\mathrm{T}}\boldsymbol{\Sigma}^{-1}(\mathbf{x}_n - \boldsymbol{\mu}_1)\big\} \\
&\qquad - \frac{1}{2}\sum_{n=1}^{N}(1-t_n)\mathrm{Tr}\big\{(\mathbf{x}_n - \boldsymbol{\mu}_2)^{\mathrm{T}}\boldsymbol{\Sigma}^{-1}(\mathbf{x}_n - \boldsymbol{\mu}_2)\big\}\bigg) \\
&= \frac{\partial}{\partial \boldsymbol{\Sigma}^{-1}}\bigg(\frac{N}{2}\ln|\boldsymbol{\Sigma}^{-1}| - \frac{1}{2}\sum_{n=1}^{N} t_n \mathrm{Tr}\big\{\boldsymbol{\Sigma}^{-1}(\mathbf{x}_n - \boldsymbol{\mu}_1)(\mathbf{x}_n - \boldsymbol{\mu}_1)^{\mathrm{T}}\big\} \\
&\qquad - \frac{1}{2}\sum_{n=1}^{N}(1-t_n)\mathrm{Tr}\big\{\boldsymbol{\Sigma}^{-1}(\mathbf{x}_n - \boldsymbol{\mu}_2)(\mathbf{x}_n - \boldsymbol{\mu}_2)^{\mathrm{T}}\big\}\bigg) \\
&= \frac{N}{2}\boldsymbol{\Sigma} - \frac{1}{2}\sum_{n=1}^{N} t_n(\mathbf{x}_n - \boldsymbol{\mu}_1)(\mathbf{x}_n - \boldsymbol{\mu}_1)^{\mathrm{T}} - \frac{1}{2}\sum_{n=1}^{N}(1-t_n)(\mathbf{x}_n - \boldsymbol{\mu}_2)(\mathbf{x}_n - \boldsymbol{\mu}_2)^{\mathrm{T}} \\
&= \mathbf{0},
\end{aligned}$$

where we used the following properties

$$\begin{aligned}\mathrm{Tr}(\mathbf{ABC}) &= \mathrm{Tr}(\mathbf{BCA}) \\ \frac{\partial}{\partial \mathbf{X}}\mathrm{Tr}(\mathbf{XA}) &= \mathbf{A}^{\mathrm{T}} \\ \frac{\partial}{\partial \mathbf{X}}\ln|\mathbf{X}| &= \mathbf{X}^{-\mathrm{T}}.\end{aligned}$$

Solving for $\boldsymbol{\Sigma}$, we obtain

$$\boldsymbol{\Sigma} = \frac{1}{N}\sum_{n=1}^{N}\big\{t_n(\mathbf{x}_n - \boldsymbol{\mu}_1)(\mathbf{x}_n - \boldsymbol{\mu}_1)^{\mathrm{T}} + (1-t_n)(\mathbf{x}_n - \boldsymbol{\mu}_2)(\mathbf{x}_n - \boldsymbol{\mu}_2)^{\mathrm{T}}\big\},$$

which is equivalent to (4.78) to (4.80).

A generalization to $K > 2$ classes can be derived using the same techniques. Consider the log likelihood function

$$\ell(\pi_k, \boldsymbol{\mu}_k, \boldsymbol{\Sigma}) = \sum_{n=1}^{N}\sum_{k=1}^{K} t_{nk}(\ln \pi_k + \ln \mathcal{N}(\mathbf{x}_n|\boldsymbol{\mu}_k, \boldsymbol{\Sigma})).$$

Setting the derivative with respect to $\boldsymbol{\Sigma}^{-1}$ to $\mathbf{0}$ while taking advantage of the above properties, that is

$$\begin{aligned}\frac{\partial}{\partial \boldsymbol{\Sigma}^{-1}}\ell(\pi_k, \boldsymbol{\mu}_k, \boldsymbol{\Sigma}) &= \frac{\partial}{\partial \boldsymbol{\Sigma}^{-1}}\bigg(-\frac{1}{2}\sum_{n=1}^{N}\sum_{k=1}^{K} t_{nk}\ln|\boldsymbol{\Sigma}| - \frac{1}{2}\sum_{n=1}^{N}\sum_{k=1}^{K} t_{nk}(\mathbf{x}_n - \boldsymbol{\mu}_k)^{\mathrm{T}}\boldsymbol{\Sigma}^{-1}(\mathbf{x}_n - \boldsymbol{\mu}_k)\bigg) \\ &= \frac{\partial}{\partial \boldsymbol{\Sigma}^{-1}}\bigg(\frac{1}{2}\sum_{n=1}^{N}\sum_{k=1}^{K} t_{nk}\ln|\boldsymbol{\Sigma}^{-1}| - \frac{1}{2}\sum_{n=1}^{N}\sum_{k=1}^{K} t_{nk}\mathrm{Tr}\big\{\boldsymbol{\Sigma}^{-1}(\mathbf{x}_n - \boldsymbol{\mu}_k)(\mathbf{x}_n - \boldsymbol{\mu}_k)^{\mathrm{T}}\big\}\bigg) \\ &= \frac{1}{2}\sum_{n=1}^{N}\sum_{k=1}^{K} t_{nk}\boldsymbol{\Sigma} - \frac{1}{2}\sum_{n=1}^{N}\sum_{k=1}^{K} t_{nk}(\mathbf{x}_n - \boldsymbol{\mu}_k)(\mathbf{x}_n - \boldsymbol{\mu}_k)^{\mathrm{T}} \\ &= \frac{N}{2}\boldsymbol{\Sigma} - \frac{1}{2}\sum_{n=1}^{N}\sum_{k=1}^{K} t_{nk}(\mathbf{x}_n - \boldsymbol{\mu}_k)(\mathbf{x}_n - \boldsymbol{\mu}_k)^{\mathrm{T}} \\ &= \mathbf{0},\end{aligned}$$

where in the second last step we used the fact

$$\sum_{n=1}^{N}\sum_{k=1}^{K} t_{nk} = N.$$

Hence, we obtain

$$\boldsymbol{\Sigma} = \frac{1}{N}\sum_{n=1}^{N}\sum_{k=1}^{K} t_{nk}(\mathbf{x}_n - \boldsymbol{\mu}_k)(\mathbf{x}_n - \boldsymbol{\mu}_k)^{\mathrm{T}},$$

which is a weighted average of the covariances of the data points assigned to each class.

4.3 Probabilistic Discriminative Models

(4.88)

This is easy to be verified using the chain rule.

$$\begin{aligned}\frac{\mathrm{d}\sigma}{\mathrm{d}a} &= \frac{\mathrm{d}}{\mathrm{d}a}\frac{1}{1+\exp(-a)} \\ &= -\frac{1}{(1+\exp(-a))^2}\cdot 1\cdot\exp(-a)\cdot(-1) \\ &= \frac{1}{1+\exp(-a)}\left(1-\frac{1}{1+\exp(-a)}\right) \\ &= \sigma(1-\sigma).\end{aligned}$$

(4.89)

This can be interpreted as *under the assumption that the probability of $\boldsymbol{\phi}_n$ belonging to class $\mathcal{C}_1$ is y_n, what is the chance of the given dataset coming into existence.*

(4.91)

$$\begin{aligned}\frac{\partial}{\partial\mathbf{w}}y_n &= \frac{\partial}{\partial\mathbf{w}}\sigma(\mathbf{w}^{\mathrm{T}}\boldsymbol{\phi}_n) \\ &= y_n(1-y_n)\boldsymbol{\phi}_n.\end{aligned}$$

Using this conclusion, we can compute the gradient of the error function with respect to $\mathbf{w}$, giving

$$\begin{aligned}\nabla_{\mathbf{w}}E(\mathbf{w}) &= -\nabla_{\mathbf{w}}\sum_{n=1}^{N}\{t_n\ln y_n + (1-t_n)\ln(1-y_n)\} \\ &= -\sum_{n=1}^{N}\left\{t_n\frac{1}{y_n}y_n(1-y_n)\boldsymbol{\phi}_n - (1-t_n)\frac{1}{1-y_n}y_n(1-y_n)\boldsymbol{\phi}_n\right\} \\ &= \sum_{n=1}^{N}(y_n-t_n)\boldsymbol{\phi}_n.\end{aligned}$$

(4.97)

For any vector $\mathbf{u}$ that is not perpendicular to all the feature vectors, since $0 < y_n < 1$, we have

$$\begin{aligned}\mathbf{u}^{\mathrm{T}}\mathbf{H}\mathbf{u} &= \mathbf{u}^{\mathrm{T}}\left(\sum_{n=1}^{N}y_n(1-y_n)\boldsymbol{\phi}_n\boldsymbol{\phi}_n^{\mathrm{T}}\right)\mathbf{u} \\ &= \sum_{n=1}^{N}y_n(1-y_n)(\mathbf{u}^{\mathrm{T}}\boldsymbol{\phi}_n\boldsymbol{\phi}_n^{\mathrm{T}}\mathbf{u}) \\ &= \sum_{n=1}^{N}y_n(1-y_n)(\boldsymbol{\phi}_n^{\mathrm{T}}\mathbf{u})^2 \\ &> 0.\end{aligned}$$

Hence, the Hessian is positive definite, which implies that the error function is convex and has a unique minimum.